

[SemAxis](#) is a method for scoring terms along a user-defined axis (e.g., positive-negative, concrete-abstract, hot-cold), which can be used for a range of empirical questions (for one example, see [Kozlowski et al. 2019](#)). In this homework, you'll implement SemAxis using word representations from Glove, and use it to explore corpus-specific conceptual associations.

Before running, install gensim with:

```
conda install gensim
```

In [5]:

```
conda install gensim
```

```
Collecting package metadata (current_repodata.json): ...working... done
Solving environment: ...working... failed with initial frozen solve. Retrying with flexible solve.
Solving environment: ...working... failed with repodata from current_repodata.json, will retry with next repodata source.
Collecting package metadata (repodata.json): ...working... done
Solving environment: ...working... done
```

```
## Package Plan ##
```

```
environment location: C:\Users\12062\Anaconda3\envs\anlp
```

```
added / updated specs:
- gensim
```

The following packages will be downloaded:

package	build	
boto3-1.18.21	pyhd3eb1b0_0	70 KB
botocore-1.21.21	pyhd3eb1b0_1	3.8 MB
bz2file-0.98	py38haa95532_1	246 KB
cython-0.29.23	py38hd77b12b_0	1.7 MB
gensim-4.0.1	py38hd77b12b_0	18.2 MB
jmespath-0.10.0	pyhd3eb1b0_0	22 KB
openssl-1.1.11	h2bbff1b_0	4.8 MB
s3transfer-0.5.0	pyhd3eb1b0_0	57 KB
smart_open-1.9.0	py_0	56 KB
Total:		29.0 MB

The following NEW packages will be INSTALLED:

boto	pkgs/main/win-64::boto-2.49.0-py38_0
boto3	pkgs/main/noarch::boto3-1.18.21-pyhd3eb1b0_0
botocore	pkgs/main/noarch::botocore-1.21.21-pyhd3eb1b0_1
bz2file	pkgs/main/win-64::bz2file-0.98-py38haa95532_1
cython	pkgs/main/win-64::cython-0.29.23-py38hd77b12b_0
gensim	pkgs/main/win-64::gensim-4.0.1-py38hd77b12b_0
jmespath	pkgs/main/noarch::jmespath-0.10.0-pyhd3eb1b0_0
s3transfer	pkgs/main/noarch::s3transfer-0.5.0-pyhd3eb1b0_0
smart_open	pkgs/main/noarch::smart_open-1.9.0-py_0

The following packages will be UPDATED:

openssl	1.1.1k-h2bbff1b_0 --> 1.1.11-h2bbff1b_0
---------	---

Downloading and Extracting Packages

Note: you may need to restart the kernel to use updated packages.

cython-0.29.23	1.7 MB		0%
cython-0.29.23	1.7 MB	3	4%
cython-0.29.23	1.7 MB	##8	29%
cython-0.29.23	1.7 MB	#####	50%
cython-0.29.23	1.7 MB	#####3	83%
cython-0.29.23	1.7 MB	#####	100%
boto3-1.18.21	70 KB		0%
boto3-1.18.21	70 KB	#####	100%
boto3-1.18.21	70 KB	#####	100%
botocore-1.21.21	3.8 MB		0%
botocore-1.21.21	3.8 MB	8	9%
botocore-1.21.21	3.8 MB	##4	24%
botocore-1.21.21	3.8 MB	#####	41%
botocore-1.21.21	3.8 MB	#####5	56%
botocore-1.21.21	3.8 MB	#####3	74%
botocore-1.21.21	3.8 MB	#####2	92%
botocore-1.21.21	3.8 MB	#####	100%
smart_open-1.9.0	56 KB		0%
smart_open-1.9.0	56 KB	##8	28%
smart_open-1.9.0	56 KB	#####	100%
smart_open-1.9.0	56 KB	#####	100%
jmespath-0.10.0	22 KB		0%
jmespath-0.10.0	22 KB	#####	100%
jmespath-0.10.0	22 KB	#####	100%
s3transfer-0.5.0	57 KB		0%
s3transfer-0.5.0	57 KB	#####	100%
s3transfer-0.5.0	57 KB	#####	100%
gensim-4.0.1	18.2 MB		0%
gensim-4.0.1	18.2 MB	1	1%
gensim-4.0.1	18.2 MB	4	5%
gensim-4.0.1	18.2 MB	8	9%
gensim-4.0.1	18.2 MB	#2	12%
gensim-4.0.1	18.2 MB	#5	16%
gensim-4.0.1	18.2 MB	#9	19%
gensim-4.0.1	18.2 MB	##2	22%
gensim-4.0.1	18.2 MB	##5	25%
gensim-4.0.1	18.2 MB	##9	29%
gensim-4.0.1	18.2 MB	###2	33%
gensim-4.0.1	18.2 MB	###5	36%
gensim-4.0.1	18.2 MB	###9	40%
gensim-4.0.1	18.2 MB	####3	43%
gensim-4.0.1	18.2 MB	####6	47%
gensim-4.0.1	18.2 MB	#####	50%
gensim-4.0.1	18.2 MB	#####3	54%
gensim-4.0.1	18.2 MB	#####7	57%
gensim-4.0.1	18.2 MB	#####	61%
gensim-4.0.1	18.2 MB	#####3	64%
gensim-4.0.1	18.2 MB	#####7	67%

gensim-4.0.1	18.2 MB	#####	71%
gensim-4.0.1	18.2 MB	#####4	74%
gensim-4.0.1	18.2 MB	#####8	78%
gensim-4.0.1	18.2 MB	#####1	81%
gensim-4.0.1	18.2 MB	#####5	85%
gensim-4.0.1	18.2 MB	#####8	89%
gensim-4.0.1	18.2 MB	#####2	93%
gensim-4.0.1	18.2 MB	#####5	96%
gensim-4.0.1	18.2 MB	#####9	99%
gensim-4.0.1	18.2 MB	#####	100%
bz2file-0.98	246 KB		0%
bz2file-0.98	246 KB	#####	100%
bz2file-0.98	246 KB	#####	100%
openssl-1.1.1l	4.8 MB		0%
openssl-1.1.1l	4.8 MB	7	8%
openssl-1.1.1l	4.8 MB	#8	18%
openssl-1.1.1l	4.8 MB	##9	30%
openssl-1.1.1l	4.8 MB	####	41%
openssl-1.1.1l	4.8 MB	#####1	51%
openssl-1.1.1l	4.8 MB	#####4	64%
openssl-1.1.1l	4.8 MB	#####5	75%
openssl-1.1.1l	4.8 MB	#####6	86%
openssl-1.1.1l	4.8 MB	#####7	97%
openssl-1.1.1l	4.8 MB	#####	100%

Preparing transaction: ...working... done
Verifying transaction: ...working... done
Executing transaction: ...working... done

```
In [6]: import re
from gensim.models import KeyedVectors
from gensim.scripts.glove2word2vec import glove2word2vec
import numpy as np
import numpy.linalg as LA
```

C:\Users\12062\Anaconda3\envs\anlp\lib\site-packages\gensim\similarities__init__.py:15: UserWarning: The gensim.similarities.levenshtein submodule is disabled, because the optional Levenshtein package <<https://pypi.org/project/python-Levenshtein/>> is unavailable. Install Levenshtein (e.g. `pip install python-Levenshtein`) to suppress this warning.
warnings.warn(msg)

In this homework, we'll be working with pre-trained word embeddings using the `gensim` library, which provides a number of functions for accessing representations for individual words and comparing them. The representations we'll use come from [Glove](#), which are trained on web data from the [Common Crawl](#) corpus.

```
In [7]: # First we have to convert the Glove format into w2v format; this creates a new file
glove_file = "../data/glove.6B.100d.100K.txt"
glove_in_w2v_format = "../data/glove.6B.100d.100K.w2v.txt"
_ = glove2word2vec(glove_file, glove_in_w2v_format)
```

<ipython-input-7-86a990e6d4aa>:4: DeprecationWarning: Call to deprecated `glove2word2vec` (KeyedVectors.load_word2vec_format(.., binary=False, no_header=True) loads GloVe text vectors.).
_ = glove2word2vec(glove_file, glove_in_w2v_format)

In [8]:

```
glove = KeyedVectors.load_word2vec_format("../data/glove.6B.100d.100K.w2v.txt", binary=
```

```
In [9]: good_vector=glove["good"]
```

Functions useful for the first question include the following:

```
In [16]: # access the representation for a single word
great_vector=glove["great"]

# use numpy to average multiple vector representations together
vecs_to_average=[good_vector, great_vector]
average=np.mean(vecs_to_average, axis=0)
# calculate the cosine similariy between two vectors
cosine_similarity=glove.cosine_similarities(good_vector, [great_vector])

print(good_vector.shape, great_vector.shape, average.shape, cosine_similarity)
```

```
(100,) (100,) (100,) [0.7592798]
```

Q1. Read the SemAxis [paper](#) and implement the SemAxis method described in sections 3.1.2 and 3.1.3. Given a set of word embeddings for positive terms $S^+ = \{v_1^+, \dots, v_n^+\}$ and embeddings for negative terms $S^- = \{v_1^-, \dots, v_n^-\}$ that define the endpoints of the axis, your output should be a single real-value score for an input word w with word representation v_w :

$$score(w)_{\mathbf{V}_{\text{axis}}} = \cos(v_w, \mathbf{V}_{\text{axis}})$$

Where:

$$\mathbf{V}^+ = \frac{1}{n} \sum_1^n v_i^+$$

$$\mathbf{V}^- = \frac{1}{m} \sum_1^m v_i^-$$

$$\mathbf{V}_{\text{axis}} = \mathbf{V}^+ - \mathbf{V}^-$$

```
In [52]: def get_semaxis_score(vectors, positive_terms=None, negative_terms=None, target_word=None):
    V_plus = []
    V_neg = []

    # your code here
    for p_word in positive_terms:
        V_plus.append(vectors[p_word])

    for n_word in negative_terms:
        V_neg.append(vectors[n_word])

    average_p = np.mean(V_plus, axis=0)
    average_n = np.mean(V_neg, axis=0)

    V_axis = np.subtract(average_p, average_n)

    # calculate the cosine similariy between two vectors
    score = vectors.cosine_similarities(vectors[target_word], [V_axis])
```

```
return score[0]
```

```
In [53]: # should be 0.342
get_semaxis_score(glove, positive_terms=["woman", "women"], negative_terms=["man", "men"])
```

```
Out[53]: 0.3424988
```

Now let's score a set of target terms along that axis

```
In [54]: def score_list_of_targets(vectors, positive_terms=None, negative_terms=None, target_words=None):
    scores=[]
    for target in target_words:
        scores.append((get_semaxis_score(vectors, positive_terms, negative_terms, target)))
    for k,v in reversed(sorted(scores)):
        print("%.3f\t%s" % (k,v))
```

```
In [55]: targets=["doctor", "nurse", "actor", "actress", "mechanic", "librarian", "architect", "chef", "cook", "magician", "mechanic"]
```

```
In [56]: score_list_of_targets(glove, positive_terms=["woman", "women"], negative_terms=["man", "men"], target_words=targets)

0.342  actress
0.294  nurse
0.219  librarian
0.106  doctor
0.024  actor
0.003  chef
-0.019  cook
-0.075  architect
-0.153  magician
-0.194  mechanic
```

Q2: Define your own concept axis by selecting a set of positive and negative terms and illustrate its utility by scoring a set of 10 target terms (as we did above).

```
In [58]: positive_terms=['patience', 'love', 'knowledge', 'perseverance', 'faith', 'hope']
negative_terms=['pride', 'greed', 'wrath', 'envy', 'lust', 'gluttony', 'sloth']
targets=['jealous', 'control', 'playful', 'competitive', 'humble', 'admire', 'frustration', 'sharing', 'mentor', 'playful', 'anger', 'jealous']

score_list_of_targets(glove, positive_terms=positive_terms, negative_terms=negative_terms, target_words=targets)

0.493  sharing
0.330  control
0.303  mentor
0.280  competitive
0.264  humble
0.199  admire
0.127  frustration
0.048  playful
0.001  anger
-0.087  jealous
```

Q3: Let's assume now that you're able to score all words in a vocabulary along several conceptual

dimensions (like the one you've defined) for a given set of word embeddings trained on a dataset. What could you do with that score? Brainstorm possible applications.

Incentivize positive word use in a community sharing environment, or parental control on a children gaming platform, banning words that are scored lower at a certain threshold, and rewardc positive attitude and word choices.